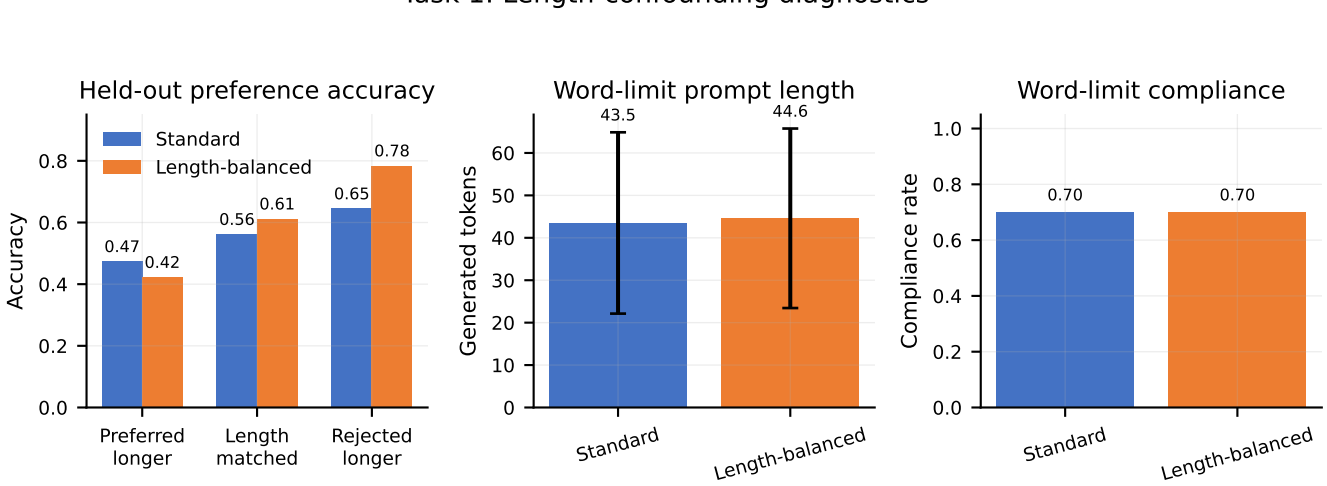
Task 1: Length-confounding diagnostics
Held-out preference accuracy
Standard
Length-balanced
Accuracy
0.0
0.2
0.4
0.6
0.8
0.47
0.42
0.56
0.61
0.65
0.78
Preferred longer
Length matched
Rejected longer
Word-limit prompt length
Generated tokens
0
10
20
30
40
50
60
43.5
44.6
Standard
Length-balanced
Word-limit compliance
Compliance rate
0.0
0.2
0.4
0.6
0.8
1.0
0.70
0.70
Standard
Length-balanced